

Teaching Neural Networks the Rules

A Selection of Solutions to Guarantee Compliance of NNs by Design

Lou Elah Süsslin
lou.suesslin@tuwien.ac.at

192.048 Seminar in Knowledge Representation and Reasoning: Neurosymbolic Artificial Intelligence
Vienna University of Technology (TU Wien)
Summer term 2025

January 19, 2026

Contents

1. Introduction
2. HMC problems and C-HMCNN(h)
3. LCMC problems and CCN(h)
4. Propositional Logic and CCN⁺
5. Summary
6. (Extra: PiShield)
7. (Extra: Hierarchy graph problems and LH-DNN)

Introduction

Neural Networks: Usefulness & Shortcomings

Usefulness

- Correlations in raw data → Accurate predictions
- Widespread success in various domains:
 - Facial Recognition
 - Financial Forecasting
 - Product Recommendation
 - Medical Sciences

Shortcomings

- Many domains require models satisfying proven requirements
- Common (Deep) NNs fail to comply with:
 - Pre-defined requirements
 - Inherent domain/problem knowledge
- Theoretical computer science and others want guarantees
- **Example:** 2D-ConvNet + 3D-RetinaNet on ROAD-R dataset
- High performance but 90% of predictions violate at least one requirement

How this has been addressed

Via Neuro-Symbolic AI methods (3 types):

- **Integrate requirements in loss function (penalize violations)**

Advantage: reduces violation incidence, improves performance

Disadvantage: cannot guarantee satisfaction

- **Post-Processing techniques (modify outputs at inference)**

Advantage: enforce compliance in final output

Disadvantage: fix-it later approach, with decay in performance

- **Integrate constraints in network topology**

("Guaranteed Compliance by Design", "Coherent-by-Construction")

Advantage: guarantee compliance, no need for post-processing

Presentation Goal

- Showcase family of neuro-symbolic frameworks
- All share common goal: make existing NNs compliant by design to given requirements or sets thereof
- Leverage existing learning capabilities of neural networks

HMC problems and C-HMCNN(h)

Introduction

Multi-label classification (MC) problem: pair $(\mathcal{A}, \mathcal{X})$ where

- \mathcal{A} : finite set of classes/labels $\{A_1, A_2, \dots\}$
- \mathcal{X} : finite dataset of labeled examples $\{(x, y)\}$
- $x \in \mathbb{R}^D$, label $y \subseteq \mathcal{A}$
- Model $m : \mathcal{A} \times \mathbb{R}^D \rightarrow [0, 1]$, outputs score $m_A(x)$
- Prediction: $m_A(x) \geq \theta \Rightarrow$ predict label A for x

Hierarchical Multi-label Classification (HMC): add constraints Π to MC problem P

- Π : finite set of hierarchy rules $A_1 \rightarrow A$ (subclass relation)
- Constraints induce an acyclic graph
- Model coherence: if $m_{A_1}(x) \geq \theta$, then $m_A(x) \geq \theta$
- Logical violation: subclass predicted but not parent
- Stronger hierarchy violation: $m_{A_1}(x) \leq m_A(x)$ must hold
- Note: logical violation \Rightarrow hierarchy violation; converse is false

HMC: Real Example

- HMC appears in domains like image classification; e.g., annotating radiological images using IRMA codes.
- Example constraint: $\text{stomach} \rightarrow \text{gastrointestinalSystem}$.
- Whenever stomach is predicted, $\text{gastrointestinalSystem}$ must also be predicted.
- Logical violation example: predicting stomach but not $\text{gastrointestinalSystem}$.

C-HMCNN(h): Two-Step Architecture

Has a two-step architecture to solve HMC problems

1. Step: Two modules

- Bottom module h : neural network of choice

Outputs: $h_{A_1}, h_A \in [0, 1]$

- Upper module: Constraint Module (CM)

- Input: h_{A_1}, h_A from h
- Output:

$$\text{CM}_{A_1} := h_{A_1}, \quad \text{CM}_A := \max(h_A, h_{A_1})$$

- For any θ : if $\text{CM}_{A_1} \geq \theta$, then $\text{CM}_A \geq \theta$
- Final prediction is given by CM output

C-HMCNN: Max Constraint Loss (CLoss)

2. Step: Novel loss function exploits hierarchy during training

- Goal: Use hierarchy constraints to guide h during training
- Approach: Modify standard BCE loss \mathcal{L} into Max Constraint Loss (CLoss)
- Decomposition:

$$\text{CLoss} = \text{CLoss}_{A_1} + \text{CLoss}_A, \quad \mathcal{L} = \mathcal{L}_{A_1} + \mathcal{L}_A$$

- Shared term:

$$\text{CLoss}_{A_1} = -y_{A_1} \ln(\text{CM}_{A_1}) - (1 - y_{A_1}) \ln(1 - \text{CM}_{A_1})$$

- Differences in parent term:

$$\text{CLoss}_A = -y_A \ln (\max(\mathbf{h}_A, \mathbf{h}_{A_1} \cdot \mathbf{y}_{A_1}))$$

$$- (1 - y_A) \ln(1 - \text{CM}_A)$$

$$\mathcal{L}_A = -y_A \ln (\text{CM}_A) - (1 - y_A) \ln(1 - \text{CM}_A)$$

- **Why:** Ground truth y_{A_1} acts as switch for subclass influence
 - If $y_{A_1} = 1$: subclass present \Rightarrow parent gets boosted
 - If $y_{A_1} = 0$: subclass absent \Rightarrow fallback to h_A

C-HMCNN(h): Problem Setup & Baselines

Setup:

- Hierarchical multi-class classification in 2D input space
- Two rectangles $R_1 < R_2$ defining classes $A_1 = R_1$, $A = R_1 \cup R_2$
- Goal: learn nested and disjoint class structures effectively

Baseline approaches:

- **f -network + post-processing f^+** : Enforces subset relation (min / max constraints), suited for $R_1 \subset R_2$
- **g -network + post-processing g^+** : Models A as union of disjoint A_1 and $A \setminus A_1$, suited for disjoint R_1, R_2

Training:

- Feedforward nets, BCE loss, Adam optimizer
- Synthetic uniform data, 20,000 epochs

C-HMCNN(h): Experimental Results

Scenario performance:

- **Subset scenario** ($R_1 \subset R_2$): C-HMCNN matches f^+ , designed for hierarchical subset coherence
- **Disjoint scenario** ($R_1 \cap R_2 = \emptyset$): C-HMCNN matches g^+ , designed for disjoint union of classes
- **Partial-overlap scenario** ($R_1 \cap R_2 \neq \emptyset, \neq R_1$): C-HMCNN outperforms both f^+ and g^+ post-processing methods

Additional observations:

- C-HMCNN learns to delegate points effectively between classes
- Produces stable predictions without explicit post-processing constraints

LCMC problems and CCN(h)

Logically Constrained Multi-Label Classification (LCMC)

LCMC extends multi-label classification by adding logical constraints Π expressed in *normal form*:

$$A_1, \dots, A_k, \neg A_{k+1}, \dots, \neg A_n \rightarrow A, \quad 0 \leq k \leq n$$

where n is the total number of literals and k the count of positive literals.

Example: “If an image contains abdomen but neither middleAbdomen nor upperAbdomen, then it must be lowerAbdomen”:

$$\text{abdomen}, \neg \text{middleAbdomen}, \neg \text{upperAbdomen} \rightarrow \text{lowerAbdomen}$$

Key assumption: All class literals are distinct within each constraint to avoid redundancy.

Crucial observation: Normal form constraints **generalize hierarchical multi-class constraints** (HMC corresponds to the special case $n = k = 1$).

Implication: LCMC uses these constraints to extend HMC’s expressiveness and logically consistent prediction notions.

Stratification of Constraint Sets II

- Set of constraints Π may contain circular definitions, causing non-unique or non-minimal solutions.
- To avoid this, we stratify Π into Π_1, \dots, Π_s .
- Stratification is a set of pairwise disjoint, non-empty subsets Π_1, \dots, Π_s of Π called strata, with

$$\Pi = \bigcup_{i=1}^s \Pi_i.$$

- Note Π_1 may be empty; thus stratification is not a strict partition.
- This ensures no circular definitions.
- Example: For

$$\Pi = \{A_1 \rightarrow A, A_2 \rightarrow A, A, \neg A_1 \rightarrow A_2\},$$

A_1 appears only in rule bodies, so

$$\Pi_1 = \emptyset.$$

A new model to tackle LCMC problems: Motivation

Example problem for a new model:

- MC problem with classes A, A_1, A_2
- Known constraints (expressed as normal form rules):
 1. $A_1 \rightarrow A$,
 2. $A_2 \rightarrow A$,
 3. $A \wedge \neg A_1 \rightarrow A_2$

These imply: $(A_1 \cup A_2) \subseteq A$

Therefore, for any model m and $x \in \mathbb{R}^D$:

$$m_A(x) \wedge \neg m_{A_1}(x) \implies m_{A_2}(x)$$

Goals for CCN:

- Leverage standard neural networks for MC
- Exploit all constraints above
- Guarantee predictions satisfy constraints
- Improve performance
- Extend HMC problems to LCMC problems

CCN(h) — Architecture

Solution proposal: CCN(h), consists of two modules + a custom loss (like C-HMCNN(h))

- **Bottom module h :** Any neural network, outputs one score per class:

$$h_A, \quad h_{A_1}, \quad h_{A_2}$$

- **Upper constraint module (CM):**

- Inputs: h_A, h_{A_1}, h_{A_2}
- Imposes constraints from normal form rules
- Outputs constrained predictions: CM_A, CM_{A_1}, CM_{A_2}

- CM enforces ($A \wedge \neg A_1 \rightarrow A_2$):

$$CM_A = \max(h_A, CM_{A_1}, CM_{A_2})$$

$$CM_{A_1} = h_{A_1}$$

$$CM_{A_2} = \max(h_{A_2}, \min(CM_A, \overline{CM}_{A_1}))$$

Custom Loss Function *CLoss* vs Standard BCE

Class	Variant	Loss Term
A	CLoss	$-y_A \ln \max(h_A, h_{A_1}y_{A_1}, h_{A_2}y_{A_2}) - \bar{y}_A \ln \overline{CM}_A$
	BCE	$-y_A \ln CM_A - \bar{y}_A \ln \overline{CM}_A$
A_1	CLoss	$-y_{A_1} \ln CM_{A_1} - \bar{y}_{A_1} \ln \overline{CM}_{A_1}$
	BCE	Same as CLoss
A_2	CLoss	$-y_{A_2} \ln \left(\max(h_{A_2}, \min(h_A y_A, h_{A_1} y_{A_1}), h_{A_1} y_{A_1}) \right)$
		$- \bar{y}_{A_2} \ln \left(1 - \max(h_{A_2}, \min(h_A y_A + y_A, h_{A_1} y_{A_1} + y_{A_1}), h_{A_1} y_{A_1} + y_{A_1}) \right)$
	BCE	$-y_{A_2} \ln CM_{A_2} - \bar{y}_{A_2} \ln \overline{CM}_{A_2}$

Key points:

- *CLoss* uses **ground truth as switches** to guide conditional prediction delegation
- Ensures training gradients point in the **correct direction**, unlike standard BCE
- Helps model learn to exploit logical constraints for better performance and consistency

CCN(h): Problem Setup & Baselines

Setup:

- Multi-label classification on 2D inputs, classes defined by regions R_1, R_2
- Goal: learn consistent label assignments under constraints (hierarchy or disjointness)
- Three scenarios:
 - **Disjoint:** $R_1 \cap R_2 = \emptyset$
 - **Subset:** $R_1 \subset R_2$
 - **Partial overlap:** $R_1 \cap R_2 \neq \emptyset$

Baselines:

- f : standard feedforward NN, trained with BCE loss
- f^+ : f + post-processing enforcing constraints (max/min operations)
- CCN(h): coherent-by-construction model integrating constraints in training (CLoss)

Training:

- Simple NN architecture (1 hidden layer, 4 neurons, tanh)
- Adam optimizer, 20,000 epochs, uniform data sampling
- Labels structured to encode logical relations $A_1 \rightarrow A$

CCN(h): Results & Insights (1/2)

Standard Neural Network (f):

- **Struggles with decision boundaries:** f struggled to learn class boundaries for A and A_2 in the synthetic experiment.
- **Performance variability:** Strongly dependent on the distribution and relative arrangement of data points (e.g., disjoint vs. overlapping R_1, R_2).
- **Requires post-processing:** Outputs often needed correction to enforce hierarchy or logical constraints.

Post-processed Network (f^+):

- **Performance decay:** Post-processing enforces constraints but can degrade performance.
- **Scenario-dependent accuracy:** In the disjoint case, f^+ wrongly predicts points in R_1 as A_2 due to constraint logic.
- **Less stable than CCN(h):** Exhibits much higher **standard deviations** in evaluation metrics across scenarios.

CCN(h): Results & Insights (2/2)

Coherent-by-Construction Network (CCN(h)):

- Integrates constraints directly into training via a Constraint Module and Loss.
- Adapts intelligently to different class relationships.
- Learns only necessary outputs, deriving others via constraints.
- Ensures coherent predictions by design — no post-processing needed.
- Provides stable and consistent performance across scenarios.

Conclusion: Embedding hard logical constraints into training yields more robust, coherent, and interpretable multi-label models compared to relying on post-processing corrections.

Integral Part of CCN(h): Procedure CompStrata(Π)


Input: Stratified set of constraints Π (Multiple can exist, logically equal)

Job 1: Resolve circular definitions through a dependency graph, with:


- Nodes = classes
- For each constraint $r \in \Pi$:
 - Positive edges from classes in $body^+(r)$ to $head(r)$
 - Negative edges from class A with $\neg A \in body^-(r)$ to $head(r)$
- Π is stratified iff the graph contains no cycles with a negative edge

Job 2: Compute stratification Π_1, \dots, Π_s and class partitions $\mathcal{A}_1, \dots, \mathcal{A}_s$ with smallest n of strata, where:


- \mathcal{A}_i pairwise disjoint, non-empty subsets, $\bigcup_i \mathcal{A}_i = \mathcal{A}$



(a) G_Π



(b) DAGs from step 1



(c) Numbers from step 2

Example:

Figure 5: Given Π as in Example 3.14, visual representation of (a) G_Π , (b) the acyclic component graph of G_Π , (c) the number assigned to each class: 1 to A_1, A_3, A_4 and 2 to A_2 .

Propositional Logic and CCN⁺

CCN⁺: Coherent-by-construction network⁺

- Extends CCN to multi-class (MC) problems
- Moves from set of constraints Π as normal rules (CCN)
- To set of **requirements** Π as propositional logic formulas (“set of clauses”) over labels \mathcal{A}

CCN⁺: Requirements as Clauses

- Clause example: $l_1 \vee l_2 \vee \dots \vee l_n$
- Each l_i is a **literal** (label $A \in \mathcal{A}$ or negation $\neg A$)
- Clause means: at least one literal true, model must predict at least one
- Assumption: each label appears once only (pos. or neg.)
- Example: $A_1 \vee A_2$
 - At least one label predicted positive
 - Equivalent to: $\neg A_1 \rightarrow A_2$ and $\neg A_2 \rightarrow A_1$
- General rule for literal l_n : $l_1 \wedge \dots \wedge \neg l_{n-1} \rightarrow l_n$
- Rules help compute label predictions bottom-up

CCN⁺: Architecture

- Input: MC problem with requirements (\mathcal{P}, Π)
- Standard neural network h produces initial output $h(x)$, possibly incoherent
- Add requirement layer **ReqL** on top of h
 - Enforces coherence with Π
 - Produces output $ReqL(x)$ using rules from Π
- Add requirements loss **ReqL loss** for training CCN⁺
 - Teaches h to respect constraints and improve predictions




Fig. 1. Basic intuition behind CCN⁺.

Basic Case Setup

- MC problem \mathcal{P} with two labels: $\mathcal{A} = \{A_1, A_2\}$
- Requirements: $\Pi = \{\neg A_1 \vee A_2\}$
- Reformulations:
 1. $A_1 \rightarrow A_2$
 2. $\neg A_2 \rightarrow \neg A_1$
- **Important:** Cannot choose both reformulations simultaneously

Reformulation 1: $A_1 \rightarrow A_2$

- This is a normal rule: compute A_2 on grounds of A_1
- Requirement layer:

$$ReqL_{A_1} = CM_{A_1} = h_{A_1}$$

$$ReqL_{A_2} = CM_{A_2} = \max(h_{A_2}, h_{A_1})$$

- Training loss uses standard BCE similar to C-HMCNN(h):

$$\begin{cases} ReqLoss_{A_1} = -y_{A_1} \log ReqL_{A_1} - (1 - y_{A_1}) \log(1 - ReqL_{A_1}) \\ ReqLoss_{A_2} = -y_{A_2} \log \max(h_{A_2}, h_{A_1} y_{A_1}) - (1 - y_{A_2}) \log(1 - ReqL_{A_2}) \end{cases}$$

Reformulation 2: $\neg A_2 \rightarrow \neg A_1$

- Need negation function f_{\neg} with strict negation property:

$$f_{\neg}(v) = 1 - v, \quad v \in [0, 1]$$

$$\text{and for } \theta \in (0, 1) : f_{\neg}(v) > \theta \Leftrightarrow v < \theta$$

- Ensures that for any label A , only one of m_A or $m_{\neg A}$ exceeds the threshold θ
- Implies that model m predicts either A or $\neg A$, and neither equals θ
- From now: **standard negation**: $f_{\neg}(v) = 1 - v$ for all $v \in [0, 1]$, entailing $\theta = 0.5$
- Now compute ReqL (for $\neg A_2 \rightarrow \neg A_1$):

$$ReqL_{A_1} = 1 - \max(1 - h_{A_1}, 1 - h_{A_2}) = \min(h_{A_1}, h_{A_2})$$

$$ReqL_{A_2} = h_{A_2}$$

- Training loss updates accordingly (inserting ReqL):

$$\begin{cases} ReqLoss_{A_1} = -y_{A_1} \log ReqL_{A_1} - (1 - y_{A_1}) \log(1 - \min(h_{A_1}, h_{A_2}(1 - y_{A_2}))) \\ ReqLoss_{A_2} = -y_{A_2} \log ReqL_{A_2} - (1 - y_{A_2}) \log(1 - ReqL_{A_2}) \end{cases}$$

Important Conclusion

- Both reformulations represent the same clause logically
- **However, choosing both simultaneously is impossible**
- The direction of dependency impacts:
 - Model expressiveness
 - Training feasibility
 - Allowed constraints (hierarchical vs. propositional)
- Next: Visualize differences between the two reformulations

Visualizing Differences – Setup

Datasets


- $\mathcal{X}_1, \mathcal{X}_2 \subset \mathbb{R}^2$
- Each point has labels from $\mathcal{A} = \{A_1, A_2\}$
- Color code:
 - Blue: A_1, A_2
 - Light blue: A_2 only
 - Red: none

Procedure


- Train each model on \mathcal{X}_1 and \mathcal{X}_2
- Compare performance with and without Requirement Layer
- Observe how different clause formulations affect decision boundaries

Models

- Some standard feedforward NN f
 - 1 hidden layer (4 tanh), sigmoid output
 - Loss: Std. BCE
- p -CCN $^+$: clause as $A_1 \rightarrow A_2$
- n -CCN $^+$: clause as $\neg A_2 \rightarrow \neg A_1$
- All models trained 20k steps with Adam, lr = 10^{-2}




(a) \mathcal{X}_1



(b) \mathcal{X}_2


Results on Dataset \mathcal{X}_1




- $p\text{-CCN}^+$ and $n\text{-CCN}^+$ outperform standard network f
- Coherent decision boundaries reflect the incorporated background knowledge
- $p\text{-CCN}^+$ yields slightly better fit for class A_2

(Color code: Blue = A_1, A_2 ; Light blue = A_2 only; Red = none)

Understanding \mathcal{X}_1 – Base Networks




(a) \mathcal{X}_1



- Key findings:
 - $p\text{-CCN}^+$: A_2 output fades in smaller rectangle, relying on A_1 via ReqL
 - $n\text{-CCN}^+$: simpler boundary for A_1 , needs ReqL to become coherent

(Color code: Blue = A_1, A_2 ; Light blue = A_2 only; Red = none)


Results on Dataset \mathcal{X}_2




- Observations:
 - $n\text{-CCN}^+ >$ all others
 - Only $n\text{-CCN}^+$ captures the "donut" (hollow) shape
 - $p\text{-CCN}^+$ and f both fail to separate inner region

(Color code: Blue = A_1, A_2 ; Light blue = A_2 only; Red = none)

Understanding \mathcal{X}_2 – Base Networks



(b) \mathcal{X}_2



- Key findings:
 - Only $n\text{-CCN}^+$ learns to suppress A_1 inside A_2 's area
 - When trained without ReqL, $n\text{-CCN}^+$ loses this capability

(Color code: Blue = A_1, A_2 ; Light blue = A_2 only; Red = none)


Model Performance

Accuracy over five runs with different dataset seeds

Model	\mathcal{X}_1 Accuracy	\mathcal{X}_2 Accuracy
$p\text{-CCN}^+$	0.984 ± 0.006	0.924 ± 0.012
$n\text{-CCN}^+$	0.963 ± 0.015	0.966 ± 0.004
Standard NN f	0.941 ± 0.004	0.915 ± 0.010

- $p\text{-CCN}^+$ performs best on \mathcal{X}_1
- $n\text{-CCN}^+$ performs best on \mathcal{X}_2
- Both CCN variants outperform standard NN f

Label Ordering in Requirement Layer



- **Color** = label, **Outline** = negated literal
- Ordering is assumed in the figure (not inferred)
- Rule head = literal with the **highest-level label**
- Example: $\neg\text{Blue} \vee \neg\text{Green} \vee \text{Yellow}$
 $\Rightarrow \text{Green} \wedge \text{Blue} \rightarrow \text{Yellow}$

- **Rule ordering:** Turning requirements in propositional logic into a set of rules impacts model performance
- **CCN⁺ formalizes this process with a function:** establishes an ordered sequence of the labels for their computation within the requirements layer
- **Level assignment (Integer to each label, not literal):**
 - Build DAG: labels = nodes; edge from rule head to body labels
 - Level 0: no incoming edges; others = longest path from any level 0 label

Result: Set of operational rules (See example for one rule)

Summary

Summarizing Coherent-by-Design Models

Findings: Where they excel:

- **Coherent-by-construction** with hard logical constraints
- Can be expressed in full **propositional logic** or stratified normal rules
- **Superior performance** by embedding domain knowledge
- **Reliably** outperforms state-of-the-art models
- **Versatility**: Use with any neural network architecture
- **Diverse applicability**: Genomics, tabular data generation, multi-label image classification
- **GPU optimization** possible (e.g., CCN+), not covered here

Coherent-by-Design Future Research

- **Label order in propositional logic** → Influences model performance
 - Systematically determine best order
- Support other **numerical relations** than linear inequalities
- Combine **soft and hard constraints**
- Add to **Explainable AI** → Leverage reasoning capabilities
 - Generate more natural explanations for network predictions

Thank you for following!

References

- Giunchiglia, E., & Lukasiewicz, T. (2021). Multi-label classification neural networks with hard logical constraints. *Journal of Artificial Intelligence Research*, 72, 759–818.
- Stoian, M. C., Tatomir, A., Lukasiewicz, T., & Giunchiglia, E. (2024). PiShield: A PyTorch Package for Learning with Requirements. *arXiv preprint arXiv:2402.18285*.
- Giunchiglia, E., Tatomir, A., Stoian, M. C., & Lukasiewicz, T. (2024). CCN+: A neuro-symbolic framework for deep learning with requirements. *International Journal of Approximate Reasoning*, 171, 109124.
- Fiaschi, L., & Cococcioni, M. (2024). Informed deep hierarchical classification: a non-standard analysis inspired approach. *arXiv preprint arXiv:2409.16956*.
- Giunchiglia, E., Stoian, M. C., & Lukasiewicz, T. (2022). Deep learning with logical constraints. *arXiv preprint arXiv:2205.00523*.

Thank you for following! Questions?

(Extra: PiShield)

PiShield: Introduction

- Python-based ML library, extends **PyTorch**
- Focus: **Requirements-driven ML**
- Goal: Help DL models meet **safety requirements** for outputs
- How: Integrate domain requirements into NN topology to ensure **compliance regardless of input**
- Adds new PyTorch layers called **Shield Layers** on top of any NN (next slide)

Implementing a Shield Layer

Training Time: Integrate after the layer computing model output

Inference Time:

- Integrate outside NN; apply to outputs to enforce compliance
- Needs:
 - Input dimension (usually network output size)
 - Path to file with requirements




Figure 1: PiShield overview.

Requirements Format

- Each requirement on a separate text line
- Expressed as either:
 - **Propositional logic formula:** in CNF, i.e., single OR-clauses combined by ANDs
 - **Linear inequalities**

Logic Formula Example

Given: Images of traffic lights from within a car

Question: Is there a traffic light? What color?

Setup: Multi-class problem, 4 labels

- y_0 — presence of traffic light
- y_1 — red light
- y_2 — yellow light
- y_3 — green light

Each image \rightarrow 4-element output vector, e.g.

[1, 0, 0, 1]

Requirements:

- $y_0 \rightarrow (y_1 \vee y_2 \vee y_3)$
- No two colors can be TRUE simultaneously

```
not y_0 or y_1 or y_2 or y_3  
not y_0 or not y_1 or not y_2  
not y_0 or not y_1 or not y_3  
not y_0 or not y_2 or not y_3
```

Linear Inequalities Example

Given: Synthetic tabular data for clinical trial

Goal: Enforce domain knowledge, e.g.,

- $\text{MaxHemoglobin} \geq \text{MinHemoglobin}$
- $\text{MaxTemp} \geq \text{MinTemp}$

Expressed as linear inequalities to incorporate constraints into learning

$$\begin{aligned}y_{-0} - y_{-1} &\geq 0 \\y_{-2} - y_{-3} &\geq 0\end{aligned}$$

Real-World Examples Overview

Domain	Scenario	Task type	Usage of Constraints
Genomics/ Bioinformatics	Functional Genomics	HCMC problem	Propositional logic reqs./CNF: Respect biological hierarchies when predicting gene functions: If a gene has specific function, it must have broader related functions (consistent with hierarchy)
Autonomous driv- ing	Road Event Detection	MC problem	Propositional logic reqs./CNF (n=243): Guarantee detected road events obey safety rules
Structured data modeling	Tabular Data Generation	Deep generative modeling	Linear inequalities: Ensure generated synthetic data respects real-world relationships between features (e.g. never min. value > max. value). Helps making data realistic

Real-World Examples Performance

Advantages:

- **Scenario 1:** Preserving hierarchical structure
- **Scenario 2:** Guarantee satisfaction of requirements + Performance boost
- **Scenario 3:** Compliance with background knowledge + Realistic data generation

Scenario	Baseline	PiShield
Functional genomics(AU(\overline{PRC}))	0.225	0.241
Autonomous driving (f-mAP)	0.288	0.303
Tabular data generation (Utility-F1)	0.430	0.458

Table 1: Aggregated performance. The best results are in **bold**.

**(Extra: Hierarchy graph problems and
LH-DNN)**

Motivation

Neural network models should solve the **hierarchy graph problem** $HC = \langle T, S, F \rangle$:

- Graph T = Tree structure organizing classes/labels
- Each data point associated with a single path S through the hierarchy
- Classification path must extend fully to the deepest possible tree level F

T organizes labels in hierarchical categories:

- Level 1 = Upper category (e.g., *clothing*)
- Level 2 = Subcategory (e.g., *bottom*)
- Level 3 = Subsubcategory (e.g., *capri pants*)

Reformulating the Problem

Reframe the HC problem in two steps:


1. Lexicographic Multi-Objective Optimization Problem (LMOP):

- LMOP = Standard optimization with multiple strictly ordered objectives
- Objectives = Correct prediction (min loss) at each hierarchy level
- Strict priority: Higher levels optimized before deeper levels

2. Theorem from Non-Standard Analysis:

- Reformulate LMOP as a single *non-standard scalar program*
- Single objective for NN: Minimize weighted sum of per-level losses, weights expressed as powers of an infinitesimal ϵ
- Enforces strict lexicographic priority by assigning ∞ -greater importance to higher-level errors


LH-DNN Model



Gradient Projection for Lexicographic Priority

What if: Derive gradient from non-standard loss?

- Problem: Updating standard NN parameters naively is not possible
- Solution: **Projector operator** P_A modifies common representation z to ensure lexicographic priorities
- Goal: Update shared params θ_s so that learning on lower-level tasks *does not degrade* performance on higher-level tasks



Performance

Tested against B-CNN:

- Convolutional neural network tailored for **hierarchical classification tasks**
- Benchmarks: Fashion catalog; small everyday objects (100, 1000 categories)

Results:

- Faster learning, even at **half the parameters**
- More **stable learning**
 - e.g., no accuracy drop across all levels when optimizing lower levels
- No accuracy compromise on **higher levels** when improving lower levels
 - B-CNN exhibits compromise

Downside:

- Projection blocks require **recomputation each batch during training → time overhead**
- Calculation depends on parameters from previous levels (constantly changing)